

Regression Analysis Project: Predicting Glucose Levels

2025-06-03

```
# Loading Libraries
suppressPackageStartupMessages({
  suppressWarnings({
    library(tidyverse)
    library(corrplot)
    library(ggplot2)
    library(reshape2)
    library(gridExtra)
    library(Metrics)
    library(car)
    library(lmtest)
    library(MASS)
  })
})
```

Introduction:

Diabetes is a chronic health issue that affects millions of people all over the world. Diabetes can lead to serious health complications if not diagnosed early. Blood Glucose is a critical indicator when it comes to monitoring diabetes. I have taken this dataset from Kaggle - <https://www.kaggle.com/datasets/mathchi/diabetes-data-set> (<https://www.kaggle.com/datasets/mathchi/diabetes-data-set>). The dataset was originally collected by National Institute of Diabetes, Digestive and Kidney Diseases as part of a study on Pima Indian women. The main aim of this project is to identify the most effective model for predicting glucose level and to evaluate its assumptions, accuracy and robustness.

Total number of observations present are 768. We have multiple predictor variables and 1 Outcome variable.

Predictor Variables are Pregnancies: Number of times pregnant. It is int datatype Glucose: Plasma glucose concentration. It is int datatype BloodPressure: Diastolic blood pressure. It is int datatype SkinThickness: Triceps skin fold thickness. It is int datatype Insulin: 2-Hour serum insulin. It is int datatype BMI: Body mass index. It is numerical datatype DiabetesPedigreeFunction: Diabetes pedigree function. It is numerical datatype Age: Age of patient. It is int datatype Target Variable is the Outcome variable which is a Binary Indicator with 0 for No Diabetes and 1 for Diabetes. It is int datatype

We are going to consider Age, BMI, Insulin, BloodPressure and DiabetesPedigreeFunction in this analysis. As they are the prime contributors to increase in Glucose levels and we are going to omit Pregnancies, SkinThickness and Outcome.

Outcome is a binary variable. It shows whether the person has diabetes or not. Since our aim is to predict Glucose's continuous response variable, binary target is not suitable as a predictor.

Pregnancies is biologically related to females and may introduce gender specific bias. So we are not considering this variable.

SkinThickness is a localized measure of body fat, whereas BMI is a broader aspect and provides a overall body composition.

```
# Loading dataset
df <- read.csv("diabetes.csv")
head(df)
```

```
##   Pregnancies Glucose BloodPressure SkinThickness Insulin   BMI
## 1           6     148           72           35        0 33.6
## 2           1      85           66           29        0 26.6
## 3           8     183           64            0        0 23.3
## 4           1      89           66           23       94 28.1
## 5           0     137           40           35     168 43.1
## 6           5     116           74            0        0 25.6
##   DiabetesPedigreeFunction Age Outcome
## 1                   0.627  50       1
## 2                   0.351  31       0
## 3                   0.672  32       1
## 4                   0.167  21       0
## 5                   2.288  33       1
## 6                   0.201  30       0
```

We are removing Outcome, Pregnancies and Skinthickness in the below part

```
# Excluding 'Outcome', 'Pregnancies', and 'SkinThickness' columns
df_noutcome <- df[, !names(df) %in% c("Outcome", "Pregnancies", "SkinThickness")]
```

Glucose range from 0 to 199, with a mean of 120.9 and a median of 117. BloodPressure range between 0 and 122, with a mean of 69.1. Insulin shows a wide spread with values between 0 to 846 with a median of 30.5. BMI averages around 31.99, with values up to 67.1. DiabetesPedigreeFunction ranges from 0.078 to 2.42 and Age spans from 21 to 81, with a mean of 33.24.

```
# Checking data structure
str(df_noutcome)
```

```
## 'data.frame':   768 obs. of  6 variables:
## $ Glucose      : int  148 85 183 89 137 116 78 115 197 125 ...
## $ BloodPressure : int  72 66 64 66 40 74 50 0 70 96 ...
## $ Insulin      : int   0 0 0 94 168 0 88 0 543 0 ...
## $ BMI          : num  33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5 0 ...
## $ DiabetesPedigreeFunction: num  0.627 0.351 0.672 0.167 2.288 ...
## $ Age          : int  50 31 32 21 33 30 26 29 53 54 ...
```

```
summary(df_noutcome)
```

```
##      Glucose      BloodPressure      Insulin      BMI
## Min.   : 0.0   Min.   : 0.00   Min.   : 0.0   Min.   : 0.00
## 1st Qu.: 99.0   1st Qu.: 62.00   1st Qu.: 0.0   1st Qu.: 27.30
## Median :117.0   Median : 72.00   Median : 30.5   Median : 32.00
## Mean   :120.9   Mean   : 69.11   Mean   : 79.8   Mean   : 31.99
## 3rd Qu.:140.2   3rd Qu.: 80.00   3rd Qu.:127.2   3rd Qu.: 36.60
## Max.   :199.0   Max.   :122.00   Max.   :846.0   Max.   : 67.10
## DiabetesPedigreeFunction      Age
## Min.   :0.0780      Min.   :21.00
## 1st Qu.:0.2437      1st Qu.:24.00
## Median :0.3725      Median :29.00
## Mean   :0.4719      Mean   :33.24
## 3rd Qu.:0.6262      3rd Qu.:41.00
## Max.   :2.4200      Max.   :81.00
```

```
sum(is.na(df_noutcome))
```

```
## [1] 0
```

```
colnames(df_noutcome)
```

```
## [1] "Glucose"      "BloodPressure"
## [3] "Insulin"      "BMI"
## [5] "DiabetesPedigreeFunction" "Age"
```

There are no missing values in the dataset.

```
# Checking missing values
colSums(is.na(df_noutcome))
```

```
##      Glucose      BloodPressure      Insulin
##      0              0              0
##      BMI DiabetesPedigreeFunction      Age
##      0              0              0
```

From Histogram, we can see that Glucose shows a slightly left skewed distribution with more values at the far higher end of the spectrum. For BloodPressure, we could see a good normal distribution. In terms of Insulin, the distribution is right skewed. For BMI, we could see a fairly normal distribution. For DiabetesPedigreeFunction, we could see a right skewed distribution and for Age also we could see a right skewed distribution.

From Boxplot graph, we can see that Glucose, BloodPressure, Insulin, BMI, DiabetesPedigreeFunction and Age have outliers, we have to deal with these outliers before working on the model.

```

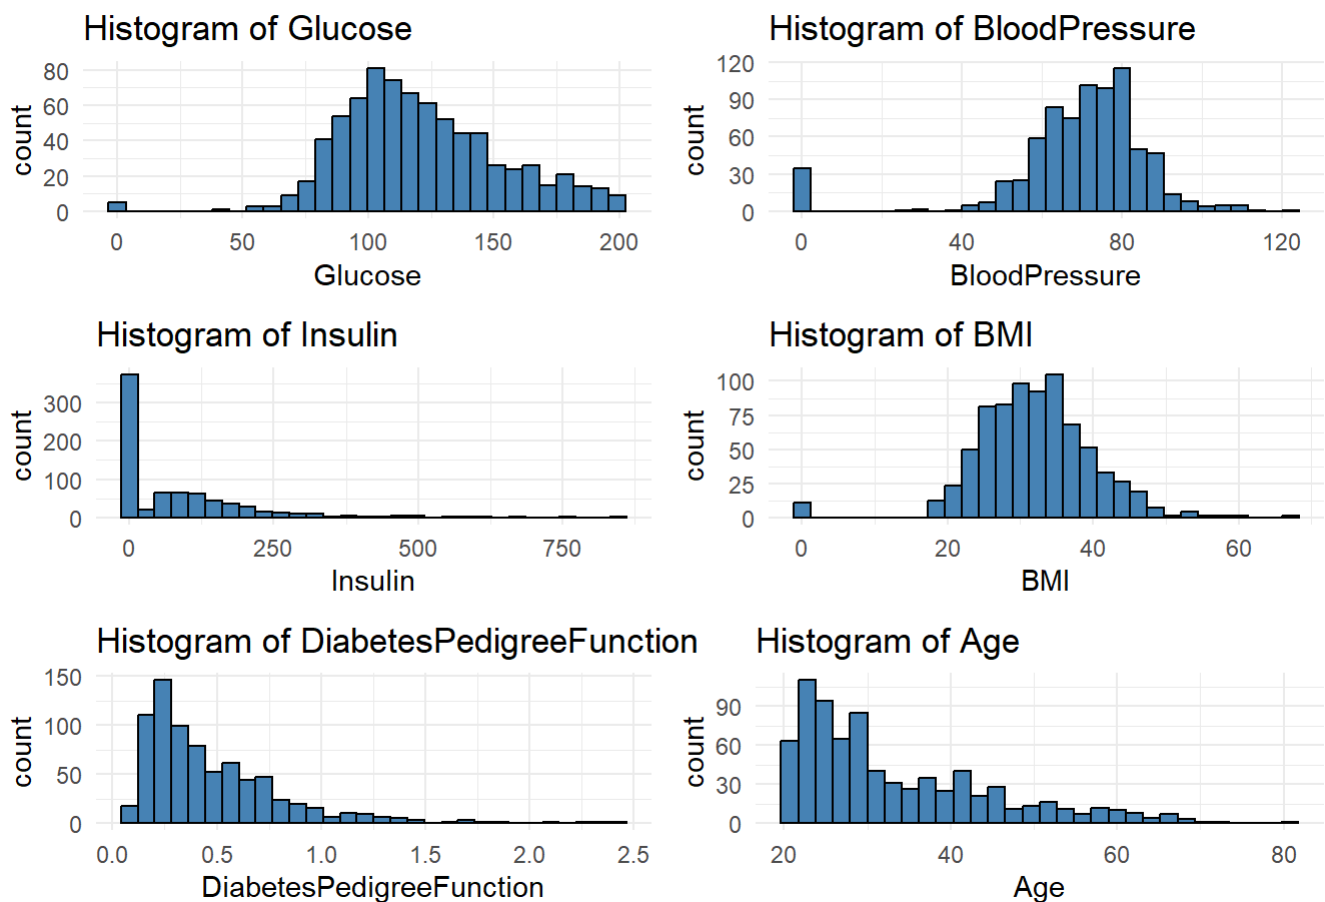
# Histogram
hist_plots <- lapply(names(df_noutcome), function(var) {
  ggplot(df_noutcome, aes(x = !!sym(var))) +
    geom_histogram(bins = 30, fill = "steelblue", color = "black") +
    theme_minimal() +
    ggtitle(paste("Histogram of", var))
})

# Boxplots
box_plots <- lapply(names(df_noutcome), function(var) {
  ggplot(df_noutcome, aes(y = !!sym(var))) +
    geom_boxplot(fill = "tomato", color = "black") +
    theme_minimal() +
    ggtitle(paste("Boxplot of", var))
})

# Displaying histogram
do.call("grid.arrange", c(hist_plots, ncol = 2, top = "Figure 1: Histogram of Predictors"))

```

Figure 1: Histogram of Predictors

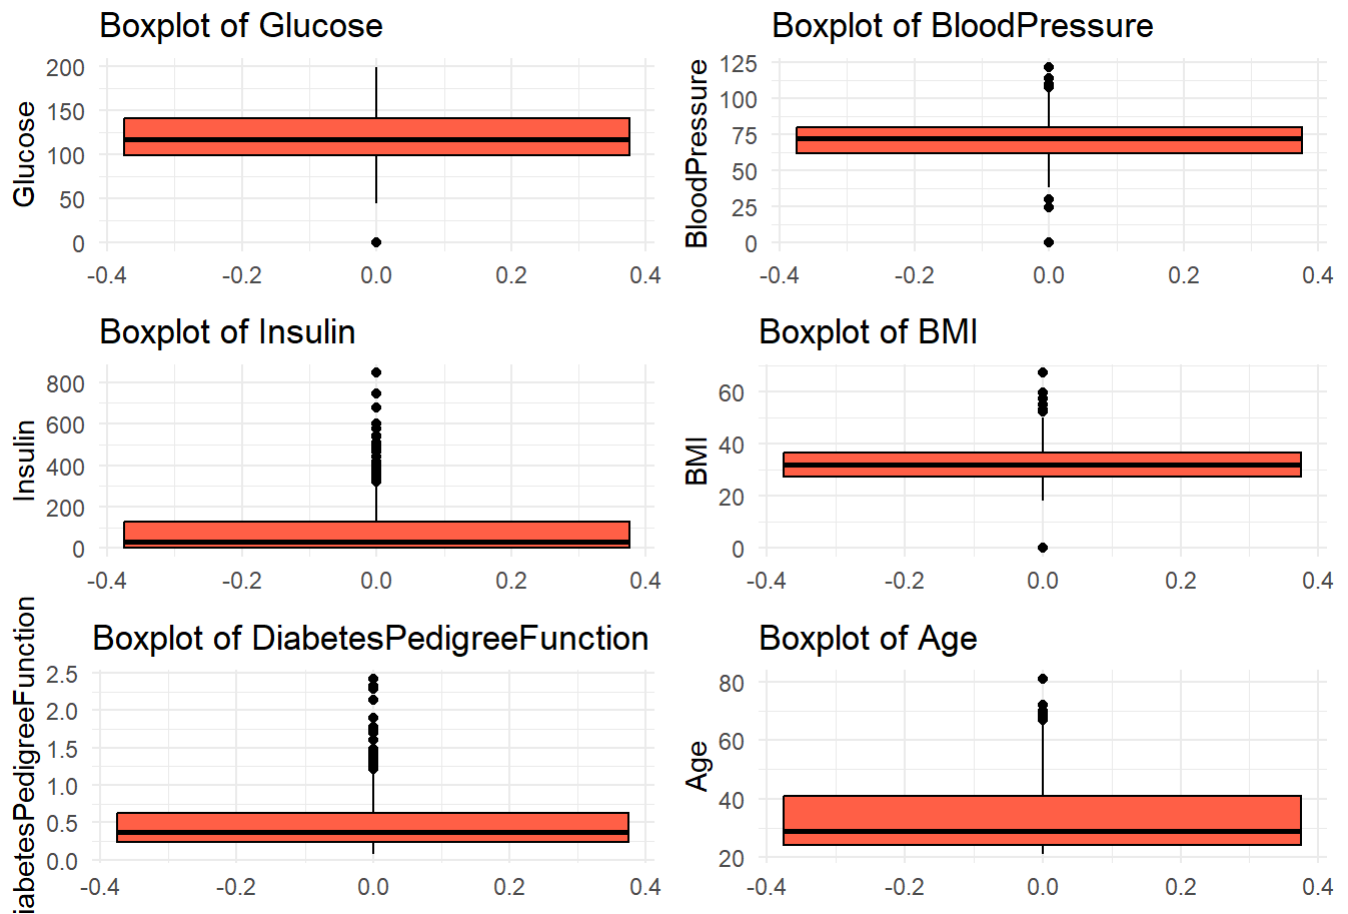


```

# Displaying boxplot
do.call("grid.arrange", c(box_plots, ncol = 2, top = "Figure 2: Boxplots of Predictors"))


```

Figure 2: Boxplots of Predictors



```
# Correlation matrix
cor_data <- df_noutcome[, c("Glucose", "Age", "BMI", "Insulin", "BloodPressure", "DiabetesPedigreeFunction")]
cor_matrix <- cor(cor_data, use = "complete.obs")

# Correlation heatmap
corrplot(cor_matrix, method = "color", addCoef.col = "black", number.cex = 0.7,
         tl.col = "black", tl.srt = 45, title = "Figure 3: Correlation Matrix", mar=c(0,0,1,0))
```

Figure 3: Correlation Matrix

In the Correlation matrix, we are focusing on Glucose vs the predictors, Insulin has 0.33 correlation with Glucose, shows moderate positive correlation. Next is Age with Glucose, which has a value of 0.26, which is weak to moderate positive correlation. BMI has a value of 0.22 which is weak positive correlation. Blood Pressure has a value of 0.15 which is very weak positive correlation. DiabetesPedigreeFunction is the weakest of all the predictors with a value of 0.14 positive correlation. Insulin is the strongest linear predictor of Glucose with Age and BMI being the next strongest predictors, Blood Pressure and DiabetesPedigreeFunction show weak correlation

We are creating a function to count outliers using the IQR method. We will apply this function to the `df_noutcome` to see the total number of Outliers present in each of the predictors.

Insulin has 34 outliers, Glucose has 5 outliers, BloodPressure has 45 which is the largest number in all of these predictors, Age has 9 outliers, BMI has 19 outliers and DiabetesPedigreeFunction has 29 outliers.

```
# Function to count outliers using IQR method
count_outliers_iqr <- function(x) {
  Q1 <- quantile(x, 0.25, na.rm = TRUE)
  Q3 <- quantile(x, 0.75, na.rm = TRUE)
  IQR <- Q3 - Q1
  lower <- Q1 - 1.5 * IQR
  upper <- Q3 + 1.5 * IQR
  sum(x < lower | x > upper, na.rm = TRUE)
}

outlier_counts <- sapply(df_noutcome, count_outliers_iqr)
outlier_counts
```

##	Glucose	BloodPressure	Insulin
##	5	45	34
##	BMI DiabetesPedigreeFunction		Age
##	19	29	9

Methods:

For Insulin, we are doing log transformation, the reason is, from the Histogram graph, we can see that the distribution is right skewed, the same applies to DiabetesPedigreeFunction. For BloodPressure, we are doing square root transformation to reduce moderate skewness and stabilize variance. For Age also its right skewed, we are doing log transformation and for BMI we are doing sqrt transformation to reduce skewness. We are not going to disturb Glucose now as it has very less outliers and in the future if we want to perform any transformations like Box-Cox, the outliers will be suppressed automatically. Also Glucose is relatively symmetric so at the moment we are not doing any changes to this response variable.

```
df_transformed <- df_noutcome

# Applying transformations
df_transformed$Insulin <- log1p(df_transformed$Insulin)
df_transformed$DiabetesPedigreeFunction <- log1p(df_transformed$DiabetesPedigreeFunction)
df_transformed$BloodPressure <- sqrt(df_transformed$BloodPressure)
df_transformed$Age <- log1p(df_transformed$Age)
df_transformed$BMI <- sqrt(df_transformed$BMI)

# Viewing summary to confirm
summary(df_transformed)
```

##	Glucose	BloodPressure	Insulin	BMI
##	Min. : 0.0	Min. : 0.000	Min. : 0.000	Min. : 0.000
##	1st Qu.: 99.0	1st Qu.: 7.874	1st Qu.: 0.000	1st Qu.: 5.225
##	Median : 117.0	Median : 8.485	Median : 3.449	Median : 5.657
##	Mean : 120.9	Mean : 8.091	Mean : 2.472	Mean : 5.584
##	3rd Qu.: 140.2	3rd Qu.: 8.944	3rd Qu.: 4.854	3rd Qu.: 6.050
##	Max. : 199.0	Max. : 11.045	Max. : 6.742	Max. : 8.191
##	DiabetesPedigreeFunction	Age		
##	Min. : 0.07511	Min. : 3.091		
##	1st Qu.: 0.21813	1st Qu.: 3.219		
##	Median : 0.31663	Median : 3.401		
##	Mean : 0.36532	Mean : 3.482		
##	3rd Qu.: 0.48628	3rd Qu.: 3.738		
##	Max. : 1.22964	Max. : 4.407		

```
sapply(df_transformed, count_outliers_iqr)
```

##	Glucose	BloodPressure	Insulin
##	5	41	0
##	BMI DiabetesPedigreeFunction		Age
##	16	13	0

After performing all the transformations, we could still see some outliers present, so the best thing we can do is construct 2 models 1.model without outliers 2.model with outliers

Compare both the model's AIC values and other aspects and see which model performs best and use it for further analysis.

We are first fitting a model with outliers.

$$E(\hat{y}) = -16.1852 + 29.2043x_1 + 5.1581x_2 + 1.9237x_3 - 0.2908x_4 + 11.4876x_5$$

```
# Model A: With outliers
model_with_outliers <- lm(Glucose ~ Age + BMI + Insulin + BloodPressure + DiabetesPedigreeFunction, data = df_transformed)
summary(model_with_outliers)
```

```
##
## Call:
## lm(formula = Glucose ~ Age + BMI + Insulin + BloodPressure +
##     DiabetesPedigreeFunction, data = df_transformed)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -129.062  -19.002   -2.789   19.402   83.515
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -16.1852    13.5529  -1.194   0.2328
## Age             29.2043     3.5822   8.153 1.46e-15 ***
## BMI              5.1581     1.2732   4.051 5.62e-05 ***
## Insulin         1.9237     0.4657   4.130 4.02e-05 ***
## BloodPressure  -0.2908     0.6061  -0.480   0.6315
## DiabetesPedigreeFunction 11.4876     5.5748   2.061   0.0397 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 29.87 on 762 degrees of freedom
## Multiple R-squared:  0.1331, Adjusted R-squared:  0.1274
## F-statistic: 23.39 on 5 and 762 DF,  p-value: < 2.2e-16
```

We are removing the outliers using a function based on the IQR rule.

```
remove_outliers <- function(x) {
  Q1 <- quantile(x, 0.25, na.rm = TRUE)
  Q3 <- quantile(x, 0.75, na.rm = TRUE)
  IQR <- Q3 - Q1
  x[x < (Q1 - 1.5 * IQR) | x > (Q3 + 1.5 * IQR)] <- NA
  return(x)
}

# Model B: Without outliers
df_clean <- as.data.frame(lapply(df_transformed, remove_outliers))
df_clean <- na.omit(df_clean)

model_without_outliers <- lm(Glucose ~ Age + BMI + Insulin + BloodPressure + DiabetesPedigreeFunction, data = df_clean)
summary(model_without_outliers)
```



```
##
## Call:
## lm(formula = Glucose ~ Age + BMI + Insulin + BloodPressure +
##     DiabetesPedigreeFunction, data = df_clean)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -75.596 -18.817  -2.267  17.560  81.438
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -50.8966    17.1304  -2.971  0.003069 **
## Age             24.5428     3.6925   6.647 6.05e-11 ***
## BMI              7.4154     2.0007   3.706 0.000227 ***
## Insulin         1.9655     0.4595   4.278 2.15e-05 ***
## BloodPressure   4.5240     1.7666   2.561 0.010651 *
## DiabetesPedigreeFunction 4.1206     6.1442   0.671 0.502667
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 28.39 on 697 degrees of freedom
## Multiple R-squared:  0.1422, Adjusted R-squared:  0.136
## F-statistic: 23.1 on 5 and 697 DF,  p-value: < 2.2e-16
```

```
pred_with <- predict(model_with_outliers)
pred_without <- predict(model_without_outliers)

# True values
actual_with <- df_transformed$Glucose
actual_without <- df_clean$Glucose

cat("With Outliers - MAE:", mae(actual_with, pred_with),
    "RMSE:", rmse(actual_with, pred_with),
    "AIC:", AIC(model_with_outliers), "\n")
```

```
## With Outliers - MAE: 23.10728 RMSE: 29.74998 AIC: 7404.874
```

```
cat("Without Outliers - MAE:", mae(actual_without, pred_without),
    "RMSE:", rmse(actual_without, pred_without),
    "AIC:", AIC(model_without_outliers), "\n")
```

```
## Without Outliers - MAE: 22.2815 RMSE: 28.26667 AIC: 6707.434
```

Model A with outliers has MAE value of 23.10, RMSE of 29.74, AIC of 7404.87, r^2 of 0.1331, Adjusted R^2 of 0.1274 whereas Model B without outliers has MAE value of 22.28, RMSE of 28.26, AIC of 6707.434, r^2 of 0.1422, Adjusted R^2 of 0.1360. MAE and RMSE are lower in Model B. Predicted value of Glucose are closer to actual values in model B. Lower RMSE suggests fewer extreme prediction errors. Lower the AIC value better is the fit of the model. Model B has lower AIC values compared to that of Model A. Model B has higher Adjusted r^2 Values which explains a slightly higher proportion of variance in Glucose. BloodPressure becomes statistically significant in Model B which wasn't the case in Model A.

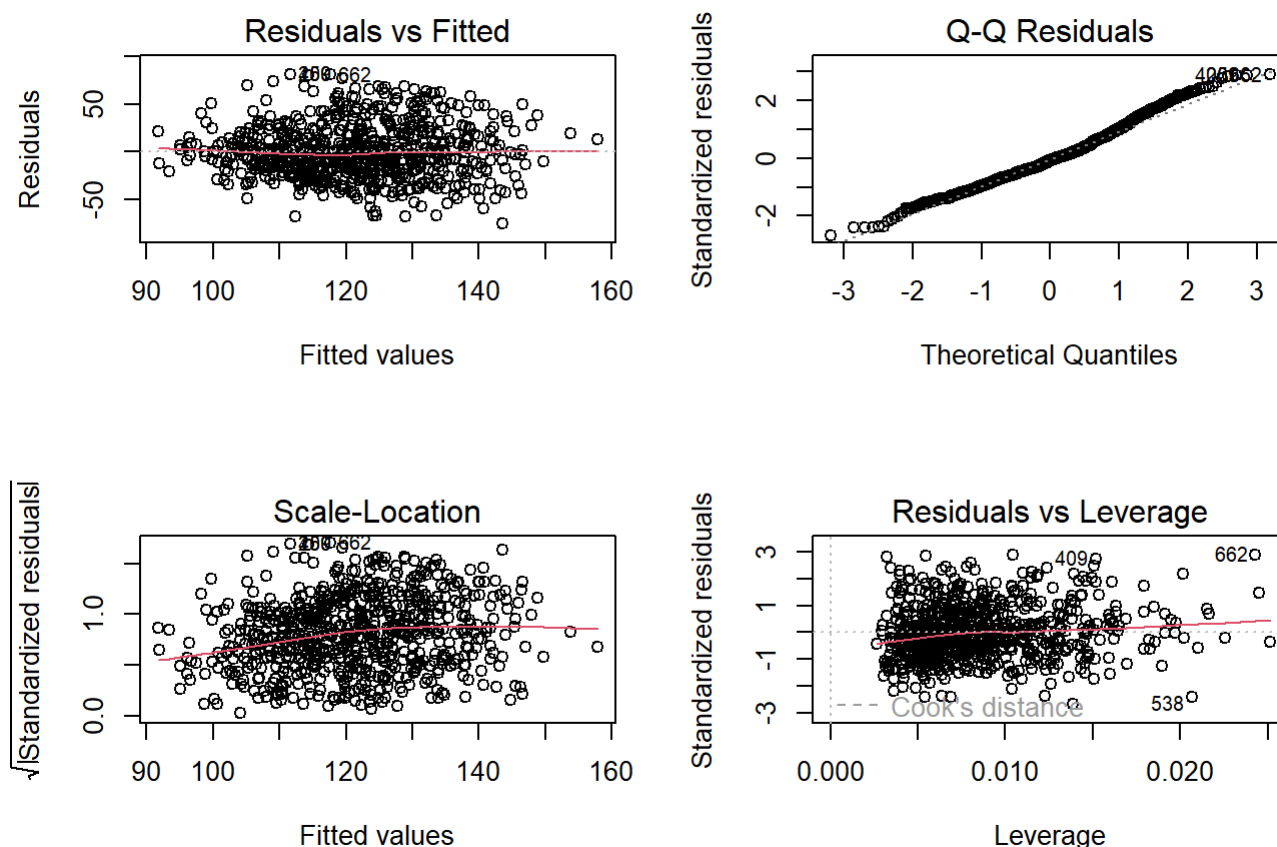
So we will go with Model B which is the model without outliers.

$$E(\hat{y}) = -50.8966 + 24.5428x_1 + 7.4154x_2 + 1.9655x_3 + 4.5240x_4 + 4.1206x_5$$


In the Residuals vs Fitted plot of the Residual Plot, we can see that all the plots are randomly scattered around the red line. In QQ-Plot all the points are on the red line, but we could see a slight variation in the top end of the tail. In Scale Location graph, we could see a slight downward trend. In Residuals vs Leverage graph, all the points are on the red line, we will further analyse the outcome of these graph with multiple tests.

```
par(mfrow=c(2,2))
plot(model_without_outliers)
mtext("Figure 4: Residual plot", side = 3, line = -2, outer = TRUE, cex = 1.2, font = 2)
```

Figure 4: Residual plot



```
# Residuals vs Fitted Plot for model_without_outliers
plot(model_without_outliers$fitted.values, model_without_outliers$residuals,
     xlab = "Fitted Values", ylab = "Residuals",
     main = "Figure 5: Residuals vs Fitted")
abline(h = 0, col = "red")
```


Figure 5: Residuals vs Fitted

```
# Standardised and Studentised Residuals
standard_res <- rstandard(model_without_outliers)
student_res <- rstudent(model_without_outliers)

# Identifying observations with high residuals
which(abs(student_res) > 2)
```

```
## 3 9 23 41 63 77 93 147 186 221 238 259 261 318 320 328 353 360 361 400
## 3 7 17 35 53 66 80 130 166 200 215 236 238 291 293 301 320 326 327 363
## 409 441 456 463 490 499 538 550 562 580 596 597 648 662 673 676 716 729 738
## 372 401 414 421 446 454 489 501 513 531 545 546 589 603 614 616 651 664 673
```

```
# Cook's Distance
cooks_d <- cooks.distance(model_without_outliers)
plot(cooks_d, type = "h",
     main = "Figure 6: Cook's Distance", ylab = "Cook's Distance")
abline(h = 4 / length(model_without_outliers$fitted.values), col = "red", lty = 2)
```


Figure 6: Cook's Distance

```
# Identifying high influence points
which(cooks_d > 4 / length(model_without_outliers$fitted.values))
```

```
## 3 9 14 23 63 93 147 155 186 221 228 238 246 259 295 320 360 363 380 400
## 3 7 10 17 53 80 130 138 166 200 206 215 223 236 269 293 326 329 344 363
## 409 441 490 499 538 550 562 580 596 597 607 648 660 662 673 716 729 748 754 764
## 372 401 446 454 489 501 513 531 545 546 552 589 601 603 614 651 664 683 689 699
```

```
#Leverage (Hat) Values
hat_vals <- hatvalues(model_without_outliers)

# Plotting Leverage values
plot(hat_vals, type = "h",
     main = "Figure 7: Leverage Values", ylab = "Leverage")
abline(h = 2 * mean(hat_vals), col = "red", lty = 2)
```

Figure 7: Leverage Values

```
# Identifying high leverage points
high_leverage <- which(hat_vals > 2 * mean(hat_vals))
print("High leverage points:")
```

```
## [1] "High leverage points:"
```

```
print(high_leverage)
```

```
## 14 34 40 44 85 103 188 244 248 295 304 457 460 465 488 520 538 550 572 577
## 10 28 34 38 72 90 168 221 225 269 277 415 418 423 444 474 489 501 523 528
## 619 662 663
## 564 603 604
```

```
# Combining all unique influential indices
final_influential_points <- unique(c(
  which(abs(student_res) > 2),
  which(cooks_d > 4 / length(model_without_outliers$fitted.values)),
  which(hat_vals > 2 * mean(hat_vals))
))

length(final_influential_points)
```

```
## [1] 69
```

Points with studentized residuals greater than 2 are considered to be outliers, we have some points above the threshold Cook's Distance shows how much the model's prediction would change if we remove a point. The points where Cook's $D > 4/n$ are considered influential points, they all will impact our regression coefficient. Leverage measures how far an observation's predictor values are from a mean of all predictor values. Points with hat value $> 2 * \text{average hat value}$ are considered high leverage. Finally in `final_influential_points`, we are combining all the observations we got from various influential points approach, we could see a total of 69 observations.

We are going to remove all the 69 observations as they have a great impact on our model.

```
# Creating a cleaned dataset by removing influential points
df_hat_removed_final <- df_clean[-final_influential_points, ]

# Fitting the final model
model_final <- lm(Glucose ~ Age + BMI + Insulin + BloodPressure + DiabetesPedigreeFunction, data = df_hat_removed_final)

# Summary of the model
summary(model_final)
```

```
##
## Call:
## lm(formula = Glucose ~ Age + BMI + Insulin + BloodPressure +
##     DiabetesPedigreeFunction, data = df_hat_removed_final)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -53.462 -17.329  -1.221  15.848  57.417
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -66.1183    15.3191  -4.316 1.84e-05 ***
## Age             26.7340     3.4082   7.844 1.88e-14 ***
## BMI             7.3335     1.8072   4.058 5.58e-05 ***
## Insulin         2.1331     0.4069   5.242 2.17e-07 ***
## BloodPressure   5.4367     1.6756   3.245 0.00124 **
## DiabetesPedigreeFunction -2.4502     5.6604  -0.433 0.66526
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 23.62 on 628 degrees of freedom
## Multiple R-squared:  0.208, Adjusted R-squared:  0.2017
## F-statistic:    33 on 5 and 628 DF,  p-value: < 2.2e-16
```

$$E(\hat{y}) = -66.1183 + 26.7340x_1 + 7.3335x_2 + 2.1331x_3 + 5.4367x_4 - 2.4502x_5$$

Hypothesis for Wald Statistics: H_0 : The predictor has no effect. H_A : Predictor has significant effect

Hypothesis for Anova

H_0 : Variable has no effect on response H_A : Variable has effect on response

From the Wald Statistics, Anova and by checking the p-value from the summary of the model, we can see that Age, BMI, Insulin, Bloodpressure have significant effect on the model because their p-value is less than 0.05 whereas DiabetesDegreeFunction has a p-value greater than 0.05 which inturn shows that it has no effect on the model. Except DiabetesDegreeFunction, rest of the predictors donot have their value in 0 for Wald Statistic,

hence we reject null hypothesis for Wald Statistics and for Anova, the predictors Age, BMI, Insulin, Bloodpressure have p-value less than 0.05, we reject H_0 for Anova and for DiabetesDegreeFunction p-value is greater than 0.05, we fail to reject H_0 . Therefore Age, BMI, Insulin, Bloodpressure are the ones which are significant predictors.

So model_final is the model without these influential points, we will check this model with other 2 models, i.e Model with outliers and model without outliers and see which model performs best in terms of low AIC value, low MAE values and lower RMSE values. After going through the results, we could see that model_final has lowest AIC, MAE and RMSE values, so we will take this model for further analysis.

```
# ANOVA
anova(model_final)
```

```
## Analysis of Variance Table
##
## Response: Glucose
##
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Age	1	51608	51608	92.5075	< 2.2e-16 ***
BMI	1	20242	20242	36.2832	2.906e-09 ***
Insulin	1	14179	14179	25.4150	6.056e-07 ***
BloodPressure	1	5904	5904	10.5829	0.001203 **
DiabetesPedigreeFunction	1	105	105	0.1874	0.665263
Residuals	628	350350	558		

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# Getting coefficients and Standard errors
coefs1 <- summary(model_final)$coefficients

# Computing Wald statistics
wald_stats <- (coefs1[, "Estimate"] / coefs1[, "Std. Error"])^2

# Computing p-values from Chi-squared distribution with df = 1
wald_pvals <- 1 - pchisq(wald_stats, df = 1)

wald_results <- data.frame(
  Estimate = coefs1[, "Estimate"],
  StdError = coefs1[, "Std. Error"],
  WaldStatistic = round(wald_stats, 3),
  p_value = round(wald_pvals, 4),
  Significance = ifelse(wald_pvals < 0.05, "Significant", "Not Significant")
)

print(wald_results)
```

```
##              Estimate StdError WaldStatistic p_value
## (Intercept)   -66.118268 15.3190532      18.629  0.0000
## Age           26.734043  3.4081815      61.530  0.0000
## BMI           7.333481  1.8071988      16.467  0.0000
## Insulin       2.133090  0.4068991      27.482  0.0000
## BloodPressure  5.436723  1.6755971      10.528  0.0012
## DiabetesPedigreeFunction -2.450180  5.6604147       0.187  0.6651
##              Significance
## (Intercept)      Significant
## Age              Significant
## BMI              Significant
## Insulin          Significant
## BloodPressure    Significant
## DiabetesPedigreeFunction Not Significant
```

```
# Function to calculate performance metrics
model_metrics <- function(model, data) {
  preds <- predict(model, data)
  actual <- data$Glucose
  mae <- mean(abs(preds - actual))
  rmse <- sqrt(mean((preds - actual)^2))
  aic <- AIC(model)
  return(c(MAE = mae, RMSE = rmse, AIC = aic))
}

# Computing metrics
metrics_with_outliers <- model_metrics(model_with_outliers, df_transformed)
metrics_without_outliers <- model_metrics(model_without_outliers, df_clean)
metrics_final <- model_metrics(model_final, df_hat_removed_final)

comparison <- rbind(
  With_Outliers = metrics_with_outliers,
  Without_Outliers = metrics_without_outliers,
  Final_Cleaned = metrics_final
)

print(round(comparison, 2))
```


```
##           MAE  RMSE   AIC
## With_Outliers  23.11 29.75 7404.87
## Without_Outliers 22.28 28.27 6707.43
## Final_Cleaned   19.11 23.51 5816.69
```

Durbin-Watson Test

Hypothesis test: H0:No autocorrelation HA:There is autocorrelation

Since DW is close to 2 and p-value is lesser than 0.05, we fail to reject H0.No correlation detected.

```
acf(model_final$residuals, main = "Figure 8: Autocorrelated Error Test")
```


Figure 8: Autocorrelated Error Test

```
dwtest(model_final)
```

```
##  
## Durbin-Watson test  
##  
## data: model_final  
## DW = 1.9637, p-value = 0.3241  
## alternative hypothesis: true autocorrelation is greater than 0
```

Breusch-Pagan Test (Heteroskedasticity)

Hypothesis test: H0:Variance is constant HA:Variance is not constant

BP = 26.814 and p-value lesser than 0.05, we reject H0. There is evidence of heteroskedasticity.

```
bptest(model_final)
```

```
##  
## studentized Breusch-Pagan test  
##  
## data: model_final  
## BP = 26.814, df = 5, p-value = 6.2e-05
```


Shapiro-Wilk Test

Hypothesis test: H0:Normal distribution is followed by residuals HA:No normality

Since p-value is lesser than 0.05, we reject H0. Residuals are not normally distributed.

```
# Q-Q Plot
qqnorm(residuals(model_final), main = "Figure 9: Q-Q Plot of Residuals")
qqline(residuals(model_final), col = "red")
```

Figure 9: Q-Q Plot of Residuals



```
# Shapiro-Wilk Test
shapiro.test(residuals(model_final))
```

```
##
##  Shapiro-Wilk normality test
##
## data:  residuals(model_final)
## W = 0.98556, p-value = 6.472e-06
```

From VIF test, we can see that no predictors have value more than 5. Hence, there is no multicollinearity present

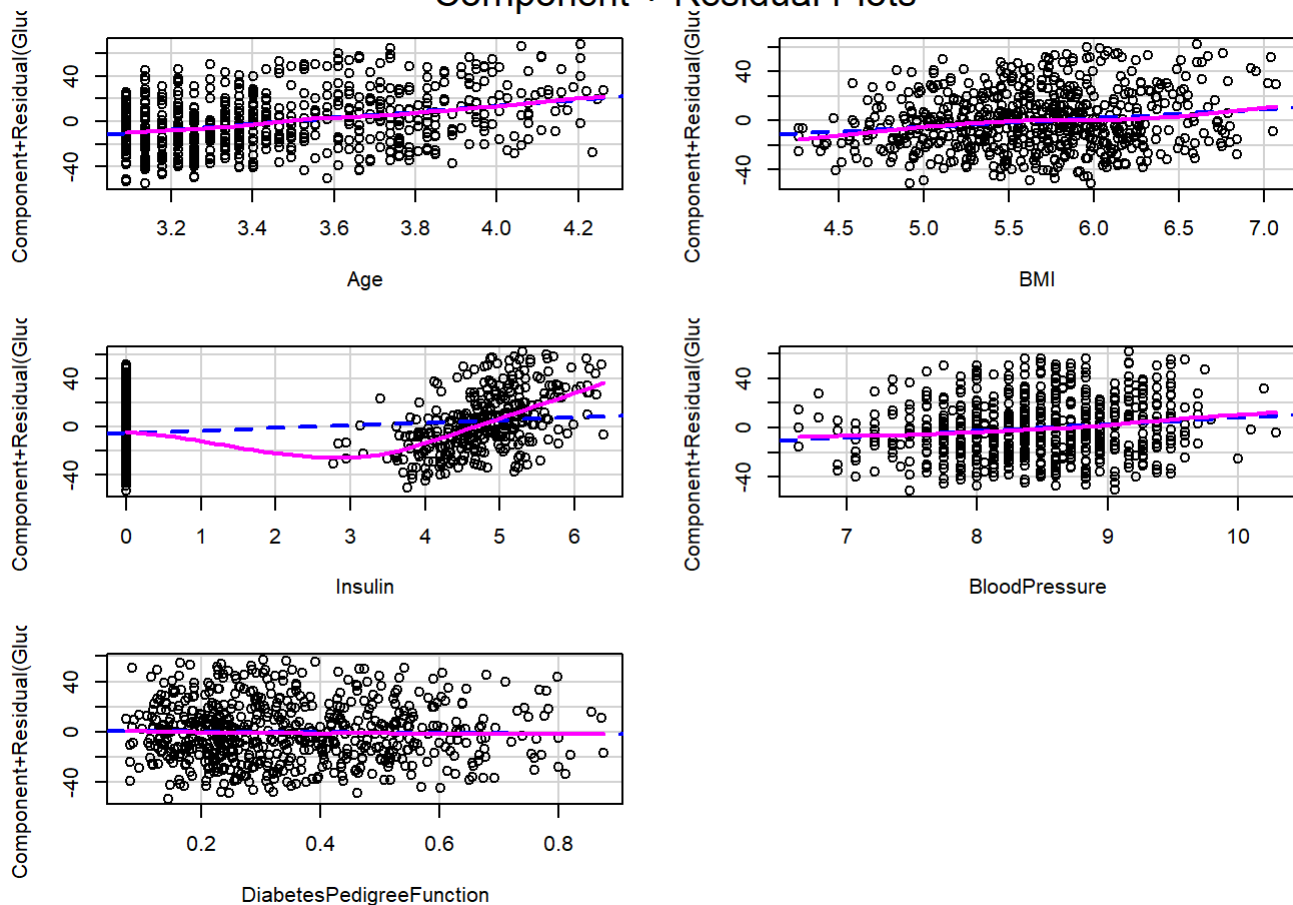
```
# VIF (Multicollinearity)
vif(model_final)
```

```
##           Age           BMI           Insulin
##      1.237229      1.146847      1.111801
## BloodPressure DiabetesPedigreeFunction
##      1.317437      1.066609
```

For Insulin alone, we could see that the pink curve is not passing through the blue line, but for remaining predictors, we can clearly see that they are all aligned properly.

```
# Linearity
car::crPlots(model_final)
mtext("Figure 10: Linearity Graph", side=3, line=3.5, cex=0.8, font=0.5)
```


Figure 10: Linearity Graph



Since we did not achieve normality and to remove heteroskedasticity, we are going to transform our model by Box-Cox transformation and see that whether we are able to achieve normality.

From Box-cox transformation, we can see that the optimal Lambda is 0.06 which is very close to 0, so we will apply log transformation to response variable Glucose and then we will fit a new model and perform all the model diagnostics.

```
# Applying Box-Cox transformation on the model
boxcox_result <- boxcox(model_final, lambda = seq(-2, 2, 0.1))
```



```
# Finding the lambda with maximum Log-Likelihood
lambda_optimal <- boxcox_result$x[which.max(boxcox_result$y)]
cat("Optimal lambda:", lambda_optimal, "\n")
```

```
## Optimal lambda: 0.06060606
```

```
# Transforming the response variable based on the optimal lambda
if (abs(lambda_optimal) < 1e-4) {
  df_transformed <- df_hat_removed_final
  df_transformed$Glucose_transformed <- log(df_transformed$Glucose)
} else {
  df_transformed <- df_hat_removed_final
  df_transformed$Glucose_transformed <- (df_transformed$Glucose^lambda_optimal - 1) / lambda_optimal
}
```

```
model_boxcox <- lm(Glucose_transformed ~ Age + BMI + Insulin + BloodPressure + DiabetesPedigr
eeFunction,
                  data = df_transformed)
summary(model_boxcox)
```

```
##
## Call:
## lm(formula = Glucose_transformed ~ Age + BMI + Insulin + BloodPressure +
##     DiabetesPedigreeFunction, data = df_transformed)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.76245 -0.17789  0.00874  0.19255  0.59476
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    3.458285    0.171850  20.124 < 2e-16 ***
## Age             0.289876    0.038233   7.582 1.23e-13 ***
## BMI             0.081488    0.020273   4.020 6.54e-05 ***
## Insulin         0.022829    0.004565   5.001 7.40e-07 ***
## BloodPressure   0.063187    0.018797   3.362 0.000822 ***
## DiabetesPedigreeFunction -0.024459    0.063499  -0.385 0.700231
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.265 on 628 degrees of freedom
## Multiple R-squared:  0.2023, Adjusted R-squared:  0.1959
## F-statistic: 31.85 on 5 and 628 DF,  p-value: < 2.2e-16
```

```
# ANOVA
anova(model_boxcox)
```

```
## Analysis of Variance Table
##
## Response: Glucose_transformed
##              Df Sum Sq Mean Sq F value    Pr(>F)
## Age             1  6.239   6.2393  88.8704 < 2.2e-16 ***
## BMI             1  2.519   2.5190  35.8797 3.535e-09 ***
## Insulin         1  1.615   1.6148  23.0013 2.024e-06 ***
## BloodPressure   1  0.797   0.7969  11.3515  0.0008 ***
## DiabetesPedigreeFunction 1  0.010   0.0104   0.1484  0.7002
## Residuals      628 44.089   0.0702
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# Getting coefficients and Standard errors
coefs2 <- summary(model_boxcox)$coefficients

# Computing Wald statistics
wald_stats1 <- (coefs2[, "Estimate"] / coefs2[, "Std. Error"])^2

# Computing p-values from Chi-squared distribution with df = 1
wald_pvals1 <- 1 - pchisq(wald_stats1, df = 1)

wald_results1 <- data.frame(
  Estimate1 = coefs2[, "Estimate"],
  StdError1 = coefs2[, "Std. Error"],
  WaldStatistic1 = round(wald_stats1, 3),
  p_value1 = round(wald_pvals1, 4),
  Significance1 = ifelse(wald_pvals1 < 0.05, "Significant", "Not Significant")
)

print(wald_results1)
```

	Estimate1	StdError1	WaldStatistic1	p_value1	Significance1
## (Intercept)	3.45828485	0.171849546	404.972	0.0000	Significant
## Age	0.28987648	0.038233071	57.484	0.0000	Significant
## BMI	0.08148847	0.020273205	16.157	0.0001	Significant
## Insulin	0.02282918	0.004564605	25.013	0.0000	Significant
## BloodPressure	0.06318721	0.018796893	11.300	0.0008	Significant
## DiabetesPedigreeFunction	-0.02445870	0.063498683	0.148	0.7001	Not Significant

$$E(\hat{y}) = 3.458285 + 0.289876x_1 + 0.081488x_2 + 0.022829x_3 + 0.063187x_4 - 0.024459x_5$$

Hypothesis for Wald Statistics: H0: The predictor has no effect. HA: Predictor has significant effect

Hypothesis for Anova

H0: Variable has no effect on response HA: Variable has effect on response

From the Wald Statistics, Anova and by checking the p-value from the summary of the model, we can see that Age, BMI, Insulin, Bloodpressure has significant effects on the model, this is same like model_final, which is the non-transformed model because they have their p-value less than 0.05 whereas DiabetesDegreeFunction has p-value greater than 0.05 and has no effect on the model. Except DiabetesDegreeFunction in this model also, rest of the predictors donot have their value in 0 for Wald Statistic, hence we reject null hypothesis for Wald Statistics and for Anova, the predictors Age, BMI, Insulin, Bloodpressure have p-value less than 0.05, we reject H0 for them and for DiabetesDegreeFunction p-value is greater than 0.05, we fail to reject H0. Therefore Age, BMI, Insulin, Bloodpressure are the ones which are significant predictors in this model too. In Box-Cox model p-value of DiabetesDegreeFunction has slightly increased compared to that of the non-transformed model.


Shapiro-Wilk Test

Hypothesis test: H0:Normal distribution is followed by residuals HA:No Normality

Since p-value is lesser than 0.05, we reject H_0 . Residuals are not normally distributed. But we can see an increase in p-value compared to non-transformed data. There is not much change in the QQ-plot where both the tails did not fit the points. But we could see a slight improvement in BC transformed model.

```
# Normality
# Q-Q Plot
qqnorm(residuals(model_boxcox), main = "Figure 11: Q-Q Plot of Residuals")
qqline(residuals(model_boxcox), col = "red")
```

Figure 11: Q-Q Plot of Residuals



```
shapiro.test(residuals(model_boxcox))
```

```
##
##  Shapiro-Wilk normality test
##
## data:  residuals(model_boxcox)
## W = 0.99129, p-value = 0.0008691
```

Breusch-Pagan Test (Heteroskedasticity)

Hypothesis test: H_0 : Variance is constant H_A : Variance is not constant

p-value is higher than 0.05, we failed to reject H_0 . There is no evidence of heteroskedasticity. After transformation heteroskedasticity is removed.

```
# Heteroscedasticity
bptest(model_boxcox)
```

```
##
## studentized Breusch-Pagan test
##
## data: model_boxcox
## BP = 6.6574, df = 5, p-value = 0.2474
```


Durbin-Watson Test

Hypothesis test: H0:No autocorrelation HA:There is autocorrelation

DW is now very close to 2 and p-value is greater than 0.05, we fail to reject H0.No correlation detected.

```
# Independence
acf(model_boxcox$residuals, main = "Figure 12: Autocorrelated Error Test")
```

Figure 12: Autocorrelated Error Test



```
dwtest(model_boxcox)
```

```
##
## Durbin-Watson test
##
## data: model_boxcox
## DW = 1.9883, p-value = 0.4418
## alternative hypothesis: true autocorrelation is greater than 0
```


No predictor is above 5, so there is no multicollinearity.


```
#Multicollinearity
vif(model_boxcox)
```

```
##                Age                BMI                Insulin
##            1.237229            1.146847            1.111801
##      BloodPressure DiabetesPedigreeFunction
##            1.317437            1.066609
```

Linearity has improved a lot for the BC-Transformed model. Insulin's value has reduced considerably and both the pink curve and the blue dotted line are very close to each other.

```
# Linearity
car::crPlots(model_boxcox)
mtext("Figure 13: Linearity Graph", side=3, line=3.5, cex=0.8, font=0.5)
```



Since Normality is still not achieved after the Box-Cox Transformation, we will fit robust model which is our second model for this analysis and see whether we are able to achieve normality for this model.

```
# Fitting robust regression on the transformed dataset
robust_model <- rlm(Glucose_transformed ~ Age + BMI + Insulin + BloodPressure + DiabetesPedigreeFunction,
                    data = df_transformed)

# Summary of the robust model
summary(robust_model)
```

```
##
## Call: rlm(formula = Glucose_transformed ~ Age + BMI + Insulin + BloodPressure +
##   DiabetesPedigreeFunction, data = df_transformed)
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-0.763921	-0.182059	0.004096	0.192065	0.593921

```
##
## Coefficients:
```

	Value	Std. Error	t value
(Intercept)	3.4426	0.1873	18.3837
Age	0.2907	0.0417	6.9785
BMI	0.0777	0.0221	3.5166
Insulin	0.0233	0.0050	4.6902
BloodPressure	0.0676	0.0205	3.3010
DiabetesPedigreeFunction	-0.0313	0.0692	-0.4524

```
##
## Residual standard error: 0.2748 on 628 degrees of freedom
```

```
# Getting coefficients matrix (no column names)
coefs3 <- summary(robust_model)$coefficients

# Assigning column names manually
colnames(coefs3) <- c("Estimate", "Std.Error", "t.value")

# Computing Wald statistics
wald_stats2 <- (coefs3[, "Estimate"] / coefs3[, "Std.Error"])^2

# Computing p-values from Chi-squared distribution with df = 1
wald_pvals2 <- 1 - pchisq(wald_stats2, df = 1)

# Compiling results
wald_results2 <- data.frame(
  Estimate2 = coefs3[, "Estimate"],
  StdError2 = coefs3[, "Std.Error"],
  WaldStatistic2 = round(wald_stats2, 3),
  p_value2 = round(wald_pvals2, 4),
  Significance2 = ifelse(wald_pvals2 < 0.05, "Significant", "Not Significant")
)

print(wald_results2)
```

##	Estimate2	StdError2	WaldStatistic2	p_value2
## (Intercept)	3.44264210	0.187265704	337.962	0.0000
## Age	0.29074285	0.041662857	48.699	0.0000
## BMI	0.07768894	0.022091859	12.367	0.0004
## Insulin	0.02332931	0.004974083	21.998	0.0000
## BloodPressure	0.06761507	0.020483112	10.897	0.0010
## DiabetesPedigreeFunction	-0.03130147	0.069194978	0.205	0.6510

##	Significance2
## (Intercept)	Significant
## Age	Significant
## BMI	Significant
## Insulin	Significant
## BloodPressure	Significant
## DiabetesPedigreeFunction	Not Significant

$$E(\hat{y}) = 3.4426 + 0.2907x_1 + 0.0777x_2 + 0.0233x_3 + 0.0676x_4 - 0.0313x_5$$

Hypothesis for Wald Statistics: H0: The predictor has no effect. HA: Predictor has significant effect


Age, BMI, Insulin and BloodPressure are significant predictors and DiabetesPedigreeFunction is not significant. We got this from the Wald Statistics. There is no Anova calculation for this model. This model's summary does not contain p-values, we manually calculated this in the Wald statistics section.

Shapiro-Wilk Test

Hypothesis test: H0: Normal distribution is followed by residuals HA: No Normality

Since p-value is lesser than 0.05, we reject H0. Residuals are not normally distributed. But we can see a increase in p-value compared to Box-Cox model. There is a slight deviation from normality. Because of the large sample size, minor normality tends to occur and in order to satisfy homoscedasticity and other measures, this violation is unlikely to affect model inference.

```
# Normality
# Q-Q Plot
qqnorm(residuals(robust_model), main = "Figure 14: Q-Q Plot of Residuals")
qqline(residuals(robust_model), col = "red")
```

Figure 14: Q-Q Plot of Residuals

```
shapiro.test(residuals(robust_model))
```

```
##
##  Shapiro-Wilk normality test
##
## data:  residuals(robust_model)
## W = 0.99153, p-value = 0.001089
```

Breusch-Pagan Test (Heteroskedasticity)

Hypothesis test: H0: Variance is constant HA: Variance is not constant

p-value is higher than 0.05, we failed to reject H0. There is no evidence of heteroskedasticity.

```
# Heteroscedasticity
bptest(robust_model)
```

```
##
##  studentized Breusch-Pagan test
##
## data:  robust_model
## BP = 6.6574, df = 5, p-value = 0.2474
```


Durbin-Watson Test

Hypothesis test: H0: No autocorrelation HA: There is autocorrelation

Since p-value is greater than 0.05, we fail to reject H0. No correlation detected.

```
# Independence
acf(robust_model$residuals, main = "Figure 15: Autocorrelated Error Test")
```

Figure 15: Autocorrelated Error Test



```
dwtest(robust_model)
```

```
##
## Durbin-Watson test
##
## data: robust_model
## DW = 1.9883, p-value = 0.4418
## alternative hypothesis: true autocorrelation is greater than 0
```

No predictor is above 5, so there is no multicollinearity.

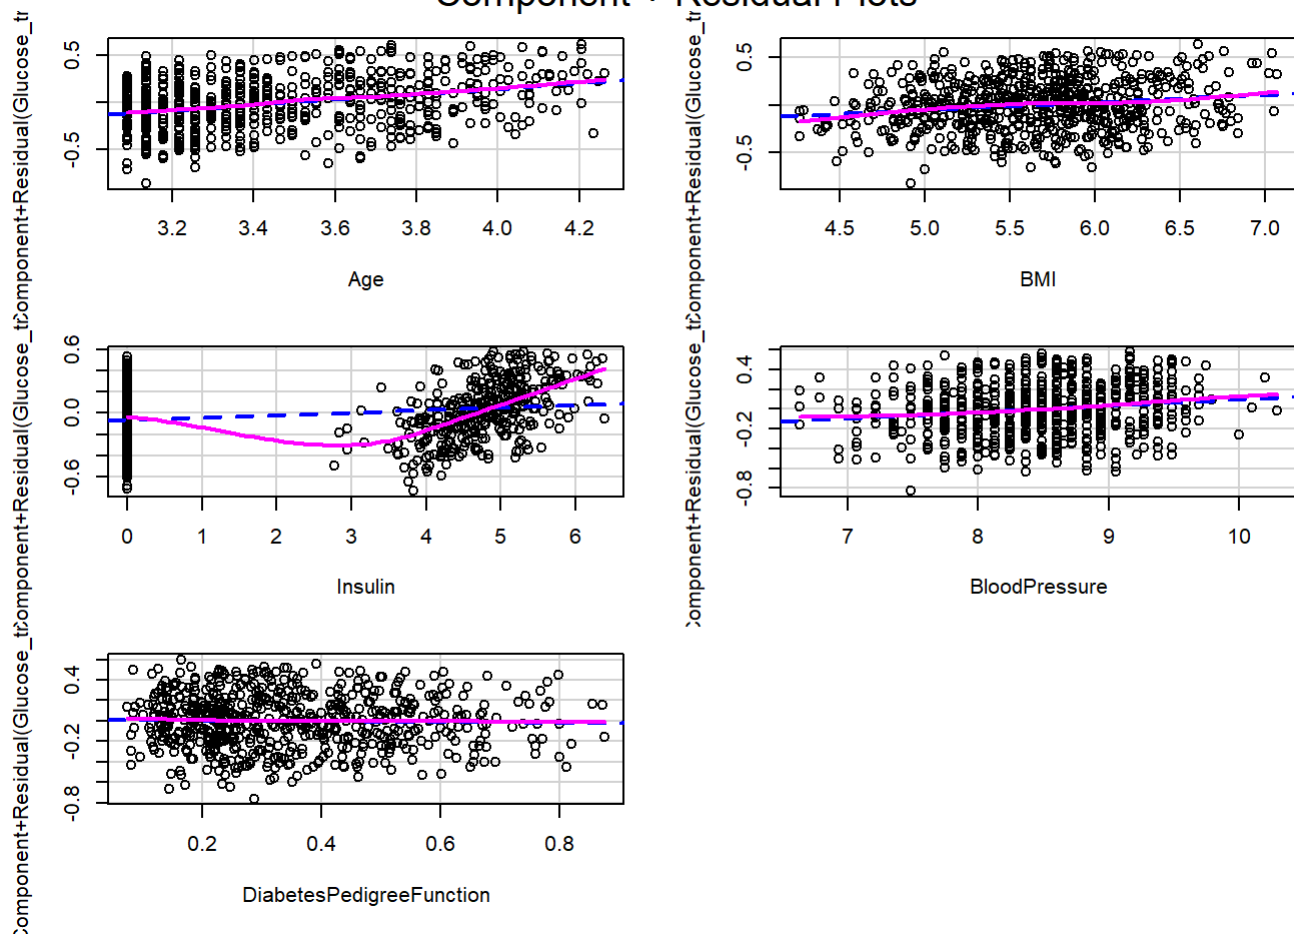
```
#Multicollinearity
vif(robust_model)
```

```
##           Age           BMI           Insulin
##      1.237229      1.146847      1.111801
## BloodPressure DiabetesPedigreeFunction
##      1.317437      1.066609
```

Linearity almost remains same as the Box-Cox model, all the predictor's both the pink curve and blue dotted line are very close to each other.

```
# Linearity
car::crPlots(robust_model)
mtext("Figure 16: Linearity Graph", side=3, line=3.5, cex=0.8, font=0.5)
```

Figure 16: Linearity Graph



```
# Making Predictions
pred_boxcox <- predict(model_boxcox)
pred_robust <- predict(robust_model)

# Actual Values
actual <- df_transformed$Glucose_transformed

# Calculate Metrics
mae_boxcox <- mae(actual, pred_boxcox)
rmse_boxcox <- rmse(actual, pred_boxcox)

mae_robust <- mae(actual, pred_robust)
rmse_robust <- rmse(actual, pred_robust)

cat("Box-Cox Model - MAE:", mae_boxcox, "RMSE:", rmse_boxcox, "\n")
```

```
## Box-Cox Model - MAE: 0.2157721 RMSE: 0.2637077
```

```
cat("Robust Model - MAE:", mae_robust, "RMSE:", rmse_robust, "\n")
```

```
## Robust Model - MAE: 0.2156808 RMSE: 0.2637381
```

```
# Residuals
resid_boxcox <- residuals(model_boxcox)
resid_robust <- residuals(robust_model)
```

```
# Histogram
hist1 <- ggplot(data.frame(resid_boxcox), aes(x = resid_boxcox)) +
  geom_histogram(bins = 30, fill = "skyblue", color = "black") +
  ggtitle("Figure 17: Residuals Histogram - BoxCox") +
  theme_minimal()

hist2 <- ggplot(data.frame(resid_robust), aes(x = resid_robust)) +
  geom_histogram(bins = 30, fill = "salmon", color = "black") +
  ggtitle("Figure 18: Residuals Histogram - Robust") +
  theme_minimal()
```

```
# QQ plots
qq1 <- ggplot(data.frame(sample = resid_boxcox), aes(sample = sample)) +
  stat_qq() + stat_qq_line() +
  ggtitle("Figure 19: QQ Plot - BoxCox") + theme_minimal()

qq2 <- ggplot(data.frame(sample = resid_robust), aes(sample = sample)) +
  stat_qq() + stat_qq_line() +
  ggtitle("Figure 20: QQ Plot - Robust") + theme_minimal()
```

```
# Fitted vs Actual
fit1 <- ggplot(df_transformed, aes(x = actual, y = pred_boxcox)) +
  geom_point(color = "blue") +
  geom_abline(slope = 1, intercept = 0, linetype = "dashed") +
  ggtitle("Figure 21: Fitted vs Actual - BoxCox") + xlab("Actual") + ylab("Predicted")

fit2 <- ggplot(df_transformed, aes(x = actual, y = pred_robust)) +
  geom_point(color = "red") +
  geom_abline(slope = 1, intercept = 0, linetype = "dashed") +
  ggtitle("Figure 22: Fitted vs Actual - Robust") + xlab("Actual") + ylab("Predicted")
```

```
grid.arrange(hist1, hist2, qq1, qq2, fit1, fit2, ncol = 2)
```

Figure 17: Residuals Histogram - BoxCox




Figure 18: Residuals Histogram - Robust




Figure 19: QQ Plot - BoxCox




Figure 20: QQ Plot - Robust




Figure 21: Fitted vs Actual - BoxCox





Figure 22: Fitted vs Actual - Robust



From the above graph, Histogram of Box-Cox is nearly symmetrical and Histogram of Robust is slightly more uniform. QQ-Plot of the models show some deviation at the tails. In the Fitted vs Actual graph, both models show similar spread. MAE and RMSE of both the models almost remain same, there is no practical difference in them.

```
# Full model with all predictors
full_model <- lm(Glucose_transformed ~ Age + BMI + Insulin + BloodPressure + DiabetesPedigree
Function,
                data = df_transformed)

# Stepwise model selection using AIC
model_stepwise <- step(full_model, direction = "both", trace = FALSE)

# Summary and AIC
summary(model_stepwise)
```



```
##
## Call:
## lm(formula = Glucose_transformed ~ Age + BMI + Insulin + BloodPressure,
##     data = df_transformed)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.76136 -0.17597  0.00782  0.19060  0.59757
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   3.458006   0.171732  20.136 < 2e-16 ***
## Age           0.288834   0.038111   7.579 1.26e-13 ***
## BMI           0.080642   0.020140   4.004 6.97e-05 ***
## Insulin       0.022467   0.004464   5.033 6.31e-07 ***
## BloodPressure 0.063320   0.018781   3.371 0.000793 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2648 on 629 degrees of freedom
## Multiple R-squared:  0.2021, Adjusted R-squared:  0.197
## F-statistic: 39.83 on 4 and 629 DF,  p-value: < 2.2e-16
```

```
AIC(model_stepwise)
```

```
## [1] 121.2288
```

```
# Predicting and Evaluating
pred_stepwise <- predict(model_stepwise, df_transformed)
actual <- df_transformed$Glucose_transformed

# Metrics
r2_stepwise <- 1 - sum((actual - pred_stepwise)^2) / sum((actual - mean(actual))^2)
rmse_stepwise <- sqrt(mean((actual - pred_stepwise)^2))
mae_stepwise <- mean(abs(actual - pred_stepwise))

cat("Stepwise Linear Model:\n")
```

```
## Stepwise Linear Model:
```

```
cat("AIC:", AIC(model_stepwise), "\n")
```

```
## AIC: 121.2288
```

```
cat("R2:", round(r2_stepwise, 4), "\nRMSE:", round(rmse_stepwise, 4), "\nMAE:", round(mae_stepwise, 4), "\n\n")
```

```
## R2: 0.2021
## RMSE: 0.2637
## MAE: 0.2158
```

From the above Stepwise model, we can see that DiabetesPedigreeFunction did not improve the model's AIC so stepwise selection removed that predictor. Remaining other predictors p-value's are lesser than 0.05. AIC of this model is 121.2288, R2 -> which means about 20.2% of the variation in Glucose is explained by this model. Adjusted R-square is 0.197. RMSE is 0.2637 and MAE is 0.2158

```
# Residuals
res <- residuals(model_stepwise)
fitted <- fitted(model_stepwise)
```

Shapiro-Wilk Test

Hypothesis test: H0:Normal distribution is followed by residuals HA:No Normality

Since p-value is lesser than 0.05, we reject H0. Residuals are not normally distributed.

```
# Shapiro-Wilk test
shapiro.test(res)
```

```
##
## Shapiro-Wilk normality test
##
## data:  res
## W = 0.99131, p-value = 0.0008803
```

Breusch-Pagan Test

Hypothesis test: H0:Variance is constant HA:Variance is not constant

p-value is higher than 0.05, we failed to reject H0. There is no evidence of heteroskedasticity.

```
# Breusch-Pagan test for heteroscedasticity
bptest(model_stepwise)
```


```
##
## studentized Breusch-Pagan test
##
## data:  model_stepwise
## BP = 5.4123, df = 4, p-value = 0.2475
```

Durbin-Watson Test

Hypothesis test: H0:No autocorrelation HA:There is autocorrelation

Since p-value is greater than 0.05, we fail to reject H0.No correlation detected.

```
# Durbin-Watson test for autocorrelation
acf(model_stepwise$residuals, main = "Figure 23: Autocorrelated Error Test")
```

Figure 23: Autocorrelated Error Test

```
durbinWatsonTest(model_stepwise)
```

```
## lag Autocorrelation D-W Statistic p-value
## 1 0.003012036 1.992283 0.938
## Alternative hypothesis: rho != 0
```

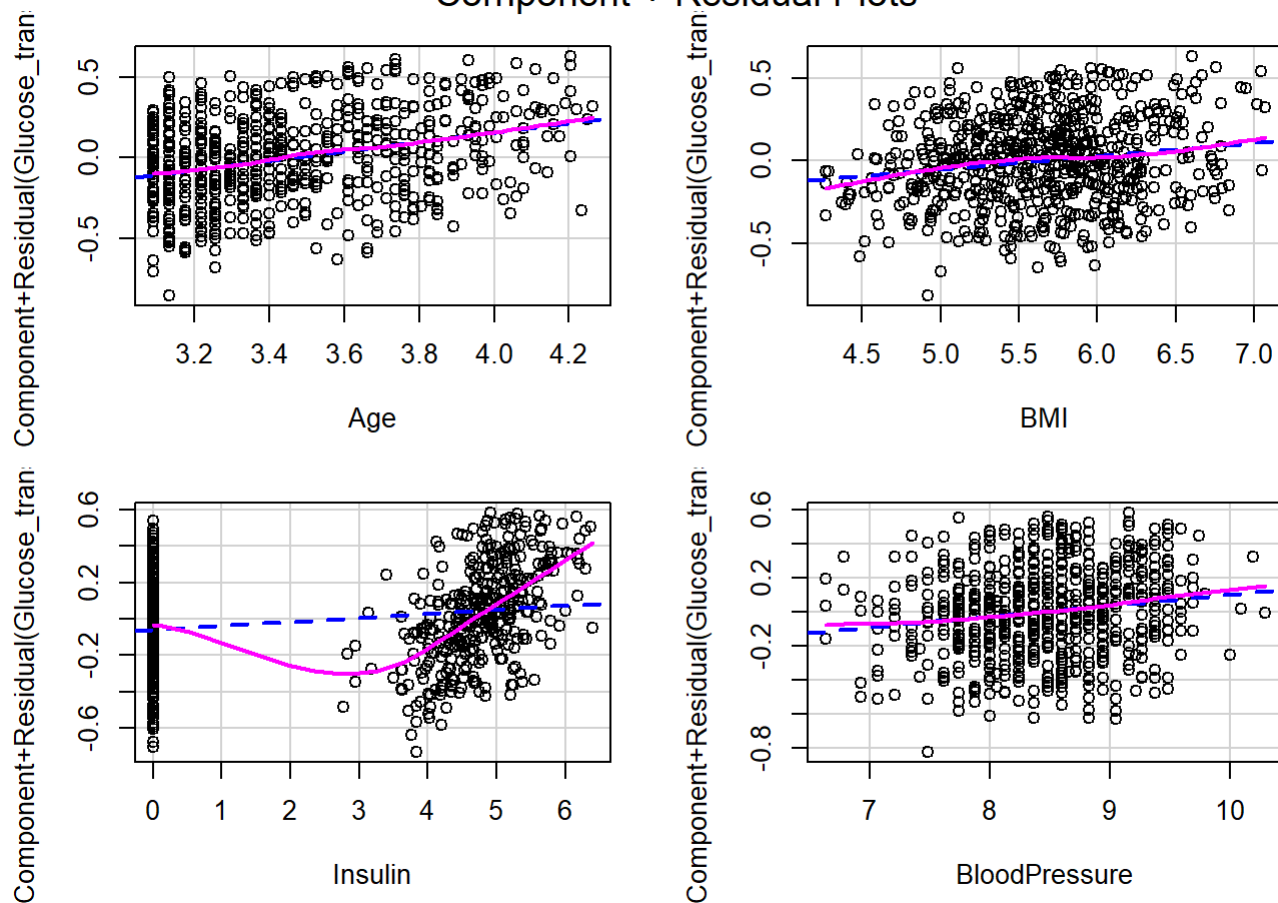
No predictor is above 5, so there is no multicollinearity.

```
# VIF for multicollinearity
vif(model_stepwise)
```

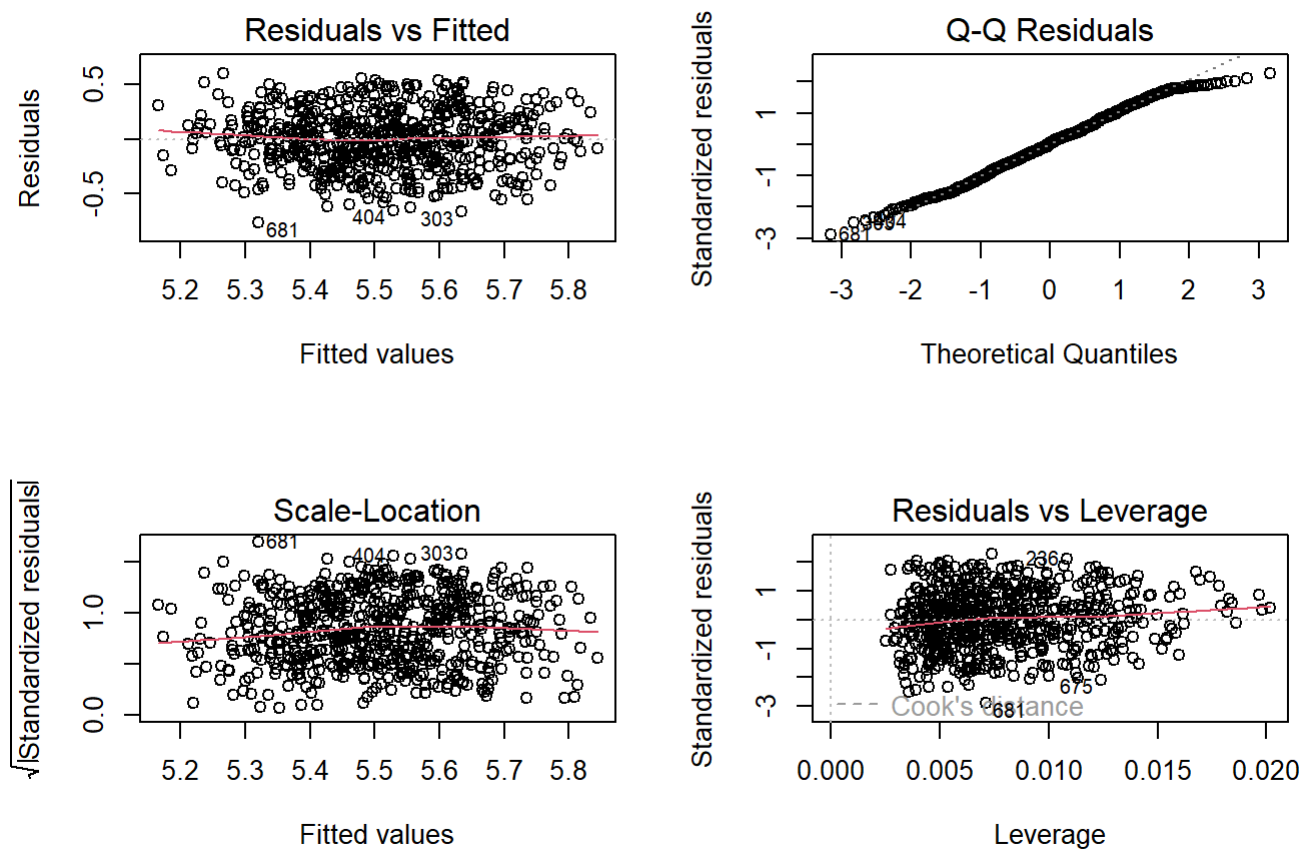
```
##      Age      BMI      Insulin BloodPressure
## 1.231026 1.133367 1.064726 1.316994
```

Linearity remains same as other 2 models, all the predictor's pink curve and blue dotted line are very close to each other.



```
# Linearity
car::crPlots(model_stepwise)
mtext("Figure 24: Linearity Graph", side=3, line=3.5, cex=0.8, font=0.5)
```

Figure 24: Linearity Graph
Component + Residual Plots

```
# Plotting diagnostics
par(mfrow = c(2, 2))
plot(model_stepwise)
mtext("Figure 25: Residual plot", side = 3, line = -2, outer = TRUE, cex = 1.2, font = 2)
```

Figure 25: Residual plot

```
# Histogram and Q-Q plot
par(mfrow = c(1, 2))
hist(res, main = "Figure 26: Histogram of Residuals", xlab = "Residuals", col = "skyblue")
qqnorm((res), main = "Figure 27: QQ Plot of Residuals")
qqline(res, col = "red")
```

Figure 26: Histogram of Residuals**Figure 27: QQ Plot of Residuals**

In Residual vs Fitted graph, all the points are well scattered around the red line. In Scale location graph, there is a slight trend of the red line. In Residuals vs Leverage graph, all the points are scattered around the red line. Histogram remains slightly symmetrical. QQ-Plot's red line tends to capture all the points in the middle but at the tail ends, it was not able to do the same.

```

# Predictions
pred_linear <- predict(model_final)
pred_robust <- predict(robust_model)
pred_stepwise <- predict(model_stepwise)
pred_boxcox <- predict(model_boxcox)

# True values
actual <- df_transformed$Glucose

# Metrics
# Linear
mae_linear <- mae(actual, pred_linear)
rmse_linear <- rmse(actual, pred_linear)
aic_linear <- AIC(model_final)

# Robust
mae_robust <- mae(actual, pred_robust)
rmse_robust <- rmse(actual, pred_robust)
aic_robust <- AIC(robust_model)

# Stepwise
mae_stepwise <- mae(actual, pred_stepwise)
rmse_stepwise <- rmse(actual, pred_stepwise)
aic_stepwise <- AIC(model_stepwise)

#BoxCox
mae_boxcox <- mae(actual, pred_boxcox)
rmse_boxcox <- rmse(actual, pred_boxcox)
aic_boxcox <- AIC(model_boxcox)

cat("\nModel Performance Comparison:\n")

```

```

##
## Model Performance Comparison:

```

```

cat("-----\n")

```

```

## -----

```

```

cat("Linear Model - MAE:", round(mae_linear, 2),
    "| RMSE:", round(rmse_linear, 2),
    "| AIC:", round(aic_linear, 2), "\n")

```

```

## Linear Model - MAE: 19.11 | RMSE: 23.51 | AIC: 5816.69

```

```

cat("Robust Model - MAE:", round(mae_robust, 2),
    "| RMSE:", round(rmse_robust, 2),
    "| AIC:", round(aic_robust, 2), "\n")

```

```

## Robust Model - MAE: 113.17 | RMSE: 116.2 | AIC: 123.23

```

```
cat("Stepwise Model      - MAE:", round(mae_stepwise, 2),  
    "| RMSE:", round(rmse_stepwise, 2),  
    "| AIC:", round(aic_stepwise, 2), "\n")
```

```
## Stepwise Model      - MAE: 113.18 | RMSE: 116.2 | AIC: 121.23
```

```
cat("BoxCox Model      - MAE:", round(mae_boxcox, 2),  
    "| RMSE:", round(rmse_boxcox, 2),  
    "| AIC:", round(aic_boxcox, 2), "\n")
```

```
## BoxCox Model      - MAE: 113.18 | RMSE: 116.2 | AIC: 123.08
```

Results and Discussion:

We have analysed various linear modelling approach to predict Glucose using the predictors Age, BMI, Insulin, BloodPressure, and DiabetesPedigreeFunction. By comparing all the 4 models MAE, RMSE and AIC, we can say that Stepwise model has the lowest AIC value which is 121.23 indicating better model parsimony whereas linear model has lowest MAE and RMSE value indicating best predictive accuracy. But I am considering AIC as the best metric among the other 2 because it balances both model fit and complexity. We should also consider parsimony while choosing a model, which the AIC provides and also it provides a more robust basis for comparison. I am going with a model which offers more efficiency and a more statistically sound result. So Stepwise Model is the best model for this dataset.

Conclusion:

After evaluating multiple linear regression models, the best model for predicting Glucose levels is Stepwise model with four significant predictors: Age, BMI, Insulin and BloodPressure. This model has the least AIC value among all the other models. While normality remains slightly violated due to the large dataset, other assumptions like independence, linearity, and multicollinearity are all satisfied with the analysis. The DiabetesPedigreeFunction had a p-value greater than 0.05 and it was thus non-significant, so it was excluded in Stepwise Model. We are considering 2 persons with Age 45 and 60, BMI is 28 and 35 and Insulin level is 90 and 150 and BloodPressure is 80 and 90. For the first patient, Predicted_Glucose value is 129.67, the value is above the normal range, typically the value should be lesser than 110. The individual is now in early diabetic range. The combination of moderate BMI and Insulin is the reason for the elevated glucose prediction. In second patient, Predicted_Glucose is 147.93, the patient is predicted to have high glucose level. Given the older age, higher BMI and higher Insulin levels, we can see a clear picture of uncontrolled diabetes. These 2 predictions reflect real world examples and clearly explain how the model can assist in early identification of people who are at risk of diabetes and help in quantifying effect of clinical measures on glucose levels.


```

# Preparing new observations
new_data <- data.frame(
  Age = log1p(c(45, 60)),
  BMI = sqrt(c(28, 35)),
  Insulin = log1p(c(90, 150)),
  BloodPressure = sqrt(c(80, 90))
)

# Prediction using stepwise model
predicted_transformed <- predict(model_stepwise, newdata = new_data)

# Applying inverse Box-Cox transformation
# Box-Cox Lambda value is
lambda <- 0.06060606

inverse_boxcox <- function(y, lambda) {
  if (abs(lambda) < 1e-4) {
    return(exp(y))
  } else {
    return((y * lambda + 1)^(1 / lambda))
  }
}

predicted_glucose <- inverse_boxcox(predicted_transformed, lambda)

# Displaying results
result_df <- cbind(
  New_Patient = c("Patient 1", "Patient 2"),
  Predicted_Glucose_Transformed = round(predicted_transformed, 4),
  Predicted_Glucose = round(predicted_glucose, 2)
)

print(result_df)

```

```

##   New_Patient Predicted_Glucose_Transformed Predicted_Glucose
## 1 "Patient 1" "5.6583"                "129.67"
## 2 "Patient 2" "5.8359"                "147.93"

```

Reference:

1. <https://www.kaggle.com/datasets/mathchi/diabetes-data-set>
(<https://www.kaggle.com/datasets/mathchi/diabetes-data-set>)
2. https://rmit.instructure.com/courses/140876/pages/week-4-after-class-3?module_item_id=7223180
(https://rmit.instructure.com/courses/140876/pages/week-4-after-class-3?module_item_id=7223180)
3. https://rmit.instructure.com/courses/140876/pages/week-5-after-class-2?module_item_id=7232820
(https://rmit.instructure.com/courses/140876/pages/week-5-after-class-2?module_item_id=7232820)
4. https://rmit.instructure.com/courses/140876/pages/week-6-after-class-2?module_item_id=7247271
(https://rmit.instructure.com/courses/140876/pages/week-6-after-class-2?module_item_id=7247271)
5. https://rmit.instructure.com/courses/140876/pages/week-7-after-class-3?module_item_id=7270551
(https://rmit.instructure.com/courses/140876/pages/week-7-after-class-3?module_item_id=7270551)

6. https://rmit.instructure.com/courses/140876/pages/week-8-after-class-2?module_item_id=7291996
(https://rmit.instructure.com/courses/140876/pages/week-8-after-class-2?module_item_id=7291996)
7. https://rmit.instructure.com/courses/140876/pages/week-9-after-class-2?module_item_id=7307061
(https://rmit.instructure.com/courses/140876/pages/week-9-after-class-2?module_item_id=7307061)
8. https://rmit.instructure.com/courses/140876/pages/week-10-after-class?module_item_id=7317487
(https://rmit.instructure.com/courses/140876/pages/week-10-after-class?module_item_id=7317487)
9. https://rmit.instructure.com/courses/140876/pages/week-11-during-class?module_item_id=7338093
(https://rmit.instructure.com/courses/140876/pages/week-11-during-class?module_item_id=7338093)

Appendix:

Loading libraries

```
suppressPackageStartupMessages({ suppressWarnings({ library(tidyverse) library(ggplot2) library(reshape2)
library(gridExtra) library(Metrics) library(car) library(lmtest) library(MASS) }) })
```

Loading dataset

```
df <- read.csv("diabetes.csv") head(df)
```

Excluding 'Outcome', 'Pregnancies', and 'SkinThickness' columns

```
df_noutcome <- df[, !names(df) %in% c("Outcome", "Pregnancies", "SkinThickness")]
```

Checking data structure

```
str(df_noutcome) summary(df_noutcome) sum(is.na(df_noutcome)) colnames(df_noutcome)
```

Checking missing values

```
colSums(is.na(df_noutcome))
```

Histogram

```
hist_plots <- lapply(names(df_noutcome), function(var) { ggplot(df_noutcome, aes(x = !!sym(var))) +
geom_histogram(bins = 30, fill = "steelblue", color = "black") + theme_minimal() + ggtitle(paste("Histogram of",
var)) })
```

Boxplots

```
box_plots <- lapply(names(df_noutcome), function(var) { ggplot(df_noutcome, aes(y = !!sym(var))) +
geom_boxplot(fill = "tomato", color = "black") + theme_minimal() + ggtitle(paste("Boxplot of", var)) })
```

Displaying histogram

```
do.call("grid.arrange", c(hist_plots, ncol = 2, top = "Figure 1: Histogram of Predictors"))
```

Displaying boxplot

```
do.call("grid.arrange", c(box_plots, ncol = 2, top = "Figure 2: Boxplots of Predictors"))
```

Correlation matrix

```
cor_data <- df_noutcome[, c("Glucose", "Age", "BMI", "Insulin", "BloodPressure", "DiabetesPedigreeFunction")]
cor_matrix <- cor(cor_data, use = "complete.obs")
```

Correlation heatmap

```
corrplot(cor_matrix, method = "color", addCoef.col = "black", number.cex = 0.7, tl.col = "black", tl.srt = 45, title =
"Figure 3: Correlation Matrix", mar=c(0,0,1,0))
```

Function to count outliers using IQR method

```
count_outliers_iqr <- function(x) { Q1 <- quantile(x, 0.25, na.rm = TRUE) Q3 <- quantile(x, 0.75, na.rm = TRUE)
IQR <- Q3 - Q1 lower <- Q1 - 1.5 * IQR upper <- Q3 + 1.5 * IQR sum(x < lower | x > upper, na.rm = TRUE) }

outlier_counts <- sapply(df_noutcome, count_outliers_iqr) outlier_counts

df_transformed <- df_noutcome
```

Applying transformations

```
df_transformed$Insulin <- -log1p(df_transformed$Insulin) df_transformed
$DiabetesPedigreeFunction <- -log1p(df_transformed$DiabetesPedigreeFunction) df_transformed
$BloodPressure <- -sqrt(df_transformed$BloodPressure) df_transformed
$Age <- -log1p(df_transformed$Age) df_transformed$BMI <- -sqrt(df_transformed$BMI)
```

Viewing summary to confirm

```
summary(df_transformed)

sapply(df_transformed, count_outliers_iqr)
```

Model A: With outliers

```
model_with_outliers <- lm(Glucose ~ Age + BMI + Insulin + BloodPressure + DiabetesPedigreeFunction, data =
df_transformed) summary(model_with_outliers)

remove_outliers <- function(x) { Q1 <- quantile(x, 0.25, na.rm = TRUE) Q3 <- quantile(x, 0.75, na.rm = TRUE)
IQR <- Q3 - Q1 x[x < (Q1 - 1.5 * IQR) | x > (Q3 + 1.5 * IQR)] <- NA return(x) }
```

Model B: Without outliers

```
df_clean <- as.data.frame(lapply(df_transformed, remove_outliers)) df_clean <- na.omit(df_clean)

model_without_outliers <- lm(Glucose ~ Age + BMI + Insulin + BloodPressure + DiabetesPedigreeFunction,
data = df_clean) summary(model_without_outliers)

pred_with <- predict(model_with_outliers) pred_without <- predict(model_without_outliers)
```

True values

```
actual_with <- df_transformed  $Glucose_{actual\_without} < -df_{clean}Glucose$ 
```

```
cat("With Outliers - MAE:", mae(actual_with, pred_with), "RMSE:", rmse(actual_with, pred_with), "AIC:",  
AIC(model_with_outliers), "")
```

```
cat("Without Outliers - MAE:", mae(actual_without, pred_without), "RMSE:", rmse(actual_without,  
pred_without), "AIC:", AIC(model_without_outliers), "")
```

```
par(mfrow=c(2,2)) plot(model_without_outliers) mtext("Figure 4: Residual plot", side = 3, line = -2, outer =  
TRUE, cex = 1.2, font = 2)
```

Residuals vs Fitted Plot for model_without_outliers

```
plot(model_without_outliers  $fitted.values$ ,  $model_{without\_outliers}residuals$ , xlab = "Fitted Values", ylab =  
"Residuals", main = "Figure 5: Residuals vs Fitted") abline(h = 0, col = "red")
```

Standardised and Studentised Residuals

```
standard_res <- rstandard(model_without_outliers) student_res <- rstudent(model_without_outliers)
```

Identifying observations with high residuals

```
which(abs(student_res) > 2)
```

Cook's Distance

```
cooks_d <- cooks.distance(model_without_outliers) plot(cooks_d, type = "h", main = "Figure 6: Cook's  
Distance", ylab = "Cook's Distance") abline(h = 4 / length(model_without_outliers  $fitted.values$ ), col = "red", lty  
= 2)
```

Identifying high influence points

```
which(cooks_d > 4 / length(model_without_outliers  $fitted.values$ ))
```

```
#Leverage (Hat) Values hat_vals <- hatvalues(model_without_outliers)
```

Plotting leverage values

```
plot(hat_vals, type = "h", main = "Figure 7: Leverage Values", ylab = "Leverage") abline(h = 2 * mean(hat_vals),  
col = "red", lty = 2)
```

Identifying high leverage points

```
high_leverage <- which(hat_vals > 2 * mean(hat_vals)) print("High leverage points:") print(high_leverage)
```

Combining all unique influential indices

```
final_influential_points <- unique(c( which(abs(student_res) > 2), which(cooks_d > 4 /  
length(model_without_outliers$fitted.values)), which(hat_vals > 2 * mean(hat_vals)) ))  
  
length(final_influential_points)
```

Creating a cleaned dataset by removing influential points

```
df_hat_removed_final <- df_clean[-final_influential_points, ]
```

Fitting the final model

```
model_final <- lm(Glucose ~ Age + BMI + Insulin + BloodPressure + DiabetesPedigreeFunction, data =  
df_hat_removed_final)
```

Summary of the model

```
summary(model_final)
```

ANOVA

```
anova(model_final)
```

Getting coefficients and Standard errors

```
coefs1 <- summary(model_final)$coefficients
```

Computing Wald statistics

```
wald_stats <- (coefs1[, "Estimate"] / coefs1[, "Std. Error"])^2
```

Computing p-values from Chi-squared distribution with $df = 1$

```
wald_pvals <- 1 - pchisq(wald_stats, df = 1)
```

```
wald_results <- data.frame( Estimate = coefs1[, "Estimate"], StdError = coefs1[, "Std. Error"], WaldStatistic =  
round(wald_stats, 3), p_value = round(wald_pvals, 4), Significance = ifelse(wald_pvals < 0.05, "Significant",  
"Not Significant") )
```

```
print(wald_results)
```

Function to calculate performance metrics

```
model_metrics <- function(model, data) { preds <- predict(model, data) actual <- data$Glucose mae <-  
mean(abs(preds - actual)) rmse <- sqrt(mean((preds - actual)^2)) aic <- AIC(model) return(c(MAE = mae,  
RMSE = rmse, AIC = aic)) }
```

Computing metrics

```
metrics_with_outliers <- model_metrics(model_with_outliers, df_transformed) metrics_without_outliers <-
model_metrics(model_without_outliers, df_clean) metrics_final <- model_metrics(model_final,
df_hat_removed_final)

comparison <- rbind( With_Outliers = metrics_with_outliers, Without_Outliers = metrics_without_outliers,
Final_Cleaned = metrics_final )

print(round(comparison, 2))

acf(model_final$residuals, main = "Figure 8: Autocorrelated Error Test") dwtest(model_final)

bptest(model_final)
```

Q-Q Plot

```
qqnorm(residuals(model_final), main = "Figure 9: Q-Q Plot of Residuals") qqline(residuals(model_final), col =
"red")
```

Shapiro-Wilk Test

```
shapiro.test(residuals(model_final))
```

VIF (Multicollinearity)

```
vif(model_final)
```

Linearity

```
car::crPlots(model_final) mtext("Figure 10: Linearity Graph", side=3, line=3.5, cex=0.8, font=0.5)
```

Applying Box-Cox transformation on the model

```
boxcox_result <- boxcox(model_final, lambda = seq(-2, 2, 0.1))
```

Finding the lambda with maximum log-likelihood

```
lambda_optimal <- boxcox_result$x[which.max(boxcox_result$y)] cat("Optimal lambda:", lambda_optimal, "")
```

Transforming the response variable based on the optimal lambda

```
if (abs(lambda_optimal) < 1e-4) { df_transformed <- df_hat_removed_final df_transformed
Glucosettransformed < -log(dfttransformedGlucose) } else { df_transformed <-
df_hat_removed_final df_transformedGlucosettransformed < -(dfttransformed
```

```
Glucose^lambda_optimal - 1) / lambda_optimal }
```

```
model_boxcox <- lm(Glucose_transformed ~ Age + BMI + Insulin + BloodPressure +
DiabetesPedigreeFunction, data = df_transformed) summary(model_boxcox)
```

ANOVA

```
anova(model_boxcox)
```

Getting coefficients and Standard errors

```
coefs2 <- summary(model_boxcox)$coefficients
```

Computing Wald statistics

```
wald_stats1 <- (coefs2[, "Estimate"] / coefs2[, "Std. Error"])^2
```

Computing p-values from Chi-squared distribution with df = 1

```
wald_pvals1 <- 1 - pchisq(wald_stats1, df = 1)
```

```
wald_results1 <- data.frame( Estimate1 = coefs2[, "Estimate"], StdError1 = coefs2[, "Std. Error"], WaldStatistic1
= round(wald_stats1, 3), p_value1 = round(wald_pvals1, 4), Significance1 = ifelse(wald_pvals1 < 0.05,
"Significant", "Not Significant") )
```

```
print(wald_results1)
```

Normality

Q-Q Plot

```
qqnorm(residuals(model_boxcox), main = "Figure 11: Q-Q Plot of Residuals") qqline(residuals(model_boxcox),
col = "red")
```

```
shapiro.test(residuals(model_boxcox))
```

Heteroscedasticity

```
bptest(model_boxcox)
```

Independence

```
acf(model_boxcox$residuals, main = "Figure 12: Autocorrelated Error Test") dwtest(model_boxcox)
```

```
#Multicollinearity vif(model_boxcox)
```

Linearity

```
car::crPlots(model_boxcox) mtext("Figure 13: Linearity Graph", side=3, line=3.5, cex=0.8, font=0.5)
```

Fitting robust regression on the transformed dataset

```
robust_model <- rlm(Glucose_transformed ~ Age + BMI + Insulin + BloodPressure +  
DiabetesPedigreeFunction, data = df_transformed)
```

Summary of the robust model

```
summary(robust_model)
```

Getting coefficients matrix (no column names)

```
coefs3 <- summary(robust_model)$coefficients
```

Assigning column names manually

```
colnames(coefs3) <- c("Estimate", "Std.Error", "t.value")
```

Computing Wald statistics

```
wald_stats2 <- (coefs3[, "Estimate"] / coefs3[, "Std.Error"])^2
```

Computing p-values from Chi-squared distribution with $df = 1$

```
wald_pvals2 <- 1 - pchisq(wald_stats2, df = 1)
```

Compiling results

```
wald_results2 <- data.frame( Estimate2 = coefs3[, "Estimate"], StdError2 = coefs3[, "Std.Error"], WaldStatistic2  
= round(wald_stats2, 3), p_value2 = round(wald_pvals2, 4), Significance2 = ifelse(wald_pvals2 < 0.05,  
"Significant", "Not Significant") )
```

```
print(wald_results2)
```

Normality

Q-Q Plot

```
qqnorm(residuals(robust_model), main = "Figure 14: Q-Q Plot of Residuals") qqline(residuals(robust_model),  
col = "red")
```

```
shapiro.test(residuals(robust_model))
```

Heteroscedasticity


```
bptest(robust_model)
```

Independence

```
acf(robust_model$residuals, main = "Figure 15: Autocorrelated Error Test") dwtest(robust_model)
```

```
#Multicollinearity vif(robust_model)
```

Linearity

```
car::crPlots(robust_model) mtext("Figure 16: Linearity Graph", side=3, line=3.5, cex=0.8, font=0.5)
```

Making Predictions

```
pred_boxcox <- predict(model_boxcox) pred_robust <- predict(robust_model)
```

Actual Values

```
actual <- df_transformed$Glucose_transformed
```

Calculate Metrics

```
mae_boxcox <- mae(actual, pred_boxcox) rmse_boxcox <- rmse(actual, pred_boxcox)
```

```
mae_robust <- mae(actual, pred_robust) rmse_robust <- rmse(actual, pred_robust)
```

```
cat("Box-Cox Model - MAE:", mae_boxcox, "RMSE:", rmse_boxcox, "") cat("Robust Model - MAE:",  
mae_robust, "RMSE:", rmse_robust, "")
```

Residuals

```
resid_boxcox <- residuals(model_boxcox) resid_robust <- residuals(robust_model)
```

Histogram

```
hist1 <- ggplot(data.frame(resid_boxcox), aes(x = resid_boxcox)) + geom_histogram(bins = 30, fill = "skyblue",  
color = "black") + ggtitle("Figure 17: Residuals Histogram - BoxCox") + theme_minimal()
```

```
hist2 <- ggplot(data.frame(resid_robust), aes(x = resid_robust)) + geom_histogram(bins = 30, fill = "salmon",  
color = "black") + ggtitle("Figure 18: Residuals Histogram - Robust") + theme_minimal()
```

QQ plots

```
qq1 <- ggplot(data.frame(sample = resid_boxcox), aes(sample = sample)) + stat_qq() + stat_qq_line() +  
ggtitle("Figure 19: QQ Plot - BoxCox") + theme_minimal()
```

```
qq2 <- ggplot(data.frame(sample = resid_robust), aes(sample = sample)) + stat_qq() + stat_qq_line() +  
ggtitle("Figure 20: QQ Plot - Robust") + theme_minimal()
```

Fitted vs Actual

```
fit1 <- ggplot(df_transformed, aes(x = actual, y = pred_boxcox)) + geom_point(color = "blue") +  
geom_abline(slope = 1, intercept = 0, linetype = "dashed") + ggtitle("Figure 21: Fitted vs Actual - BoxCox") +  
xlab("Actual") + ylab("Predicted")  
  
fit2 <- ggplot(df_transformed, aes(x = actual, y = pred_robust)) + geom_point(color = "red") +  
geom_abline(slope = 1, intercept = 0, linetype = "dashed") + ggtitle("Figure 22: Fitted vs Actual - Robust") +  
xlab("Actual") + ylab("Predicted")  
  
grid.arrange(hist1, hist2, qq1, qq2, fit1, fit2, ncol = 2)
```

Full model with all predictors

```
full_model <- lm(Glucose_transformed ~ Age + BMI + Insulin + BloodPressure + DiabetesPedigreeFunction,  
data = df_transformed)
```

Stepwise model selection using AIC

```
model_stepwise <- step(full_model, direction = "both", trace = FALSE)
```

Summary and AIC

```
summary(model_stepwise) AIC(model_stepwise)
```

Predicting and Evaluating

```
pred_stepwise <- predict(model_stepwise, df_transformed) actual <- df_transformed$Glucose_transformed
```

Metrics

```
r2_stepwise <- 1 - sum((actual - pred_stepwise)^2) / sum((actual - mean(actual))^2) rmse_stepwise <-  
sqrt(mean((actual - pred_stepwise)^2)) mae_stepwise <- mean(abs(actual - pred_stepwise))  
  
cat("Stepwise Linear Model:") cat("AIC:", AIC(model_stepwise), "") cat("R2:", round(r2_stepwise, 4), ".",  
round(rmse_stepwise, 4), ".", round(mae_stepwise, 4), "")
```

Residuals

```
res <- residuals(model_stepwise) fitted <- fitted(model_stepwise)
```

Shapiro-Wilk test

```
shapiro.test(res)
```

Breusch-Pagan test for heteroscedasticity

```
bptest(model_stepwise)
```

Durbin-Watson test for autocorrelation

```
acf(model_stepwise$residuals, main = "Figure 23: Autocorrelated Error Test")
durbinWatsonTest(model_stepwise)
```

VIF for multicollinearity

```
vif(model_stepwise)
```

Linearity

```
car::crPlots(model_stepwise) mtext("Figure 24: Linearity Graph", side=3, line=3.5, cex=0.8, font=0.5)
```

Plotting diagnostics

```
par(mfrow = c(2, 2)) plot(model_stepwise) mtext("Figure 25: Residual plot", side = 3, line = -2, outer = TRUE,
cex = 1.2, font = 2)
```

Histogram and Q-Q plot

```
par(mfrow = c(1, 2)) hist(res, main = "Figure 26: Histogram of Residuals", xlab = "Residuals", col = "skyblue")
qqnorm((res), main = "Figure 27: QQ Plot of Residuals") qqline(res, col = "red")
```

Predictions

```
pred_linear <- predict(model_final) pred_robust <- predict(robust_model) pred_stepwise <-
predict(model_stepwise) pred_boxcox <- predict(model_boxcox)
```

True values

```
actual <- df_transformed$Glucose
```

Metrics

Linear

```
mae_linear <- mae(actual, pred_linear) rmse_linear <- rmse(actual, pred_linear) aic_linear <- AIC(model_final)
```

Robust

```
mae_robust <- mae(actual, pred_robust) rmse_robust <- rmse(actual, pred_robust) aic_robust <-
AIC(robust_model)
```

Stepwise

```
mae_stepwise <- mae(actual, pred_stepwise) rmse_stepwise <- rmse(actual, pred_stepwise) aic_stepwise <-
AIC(model_stepwise)
```

```
#BoxCox mae_boxcox <- mae(actual, pred_boxcox) rmse_boxcox <- rmse(actual, pred_boxcox) aic_boxcox <-
AIC(model_boxcox)
```

```
cat("Performance Comparison:") cat("_____") cat("Linear Model - MAE:", round(mae_linear,
2), "| RMSE:", round(rmse_linear, 2), "| AIC:", round(aic_linear, 2), ")") cat("Robust Model - MAE:",
round(mae_robust, 2), "| RMSE:", round(rmse_robust, 2), "| AIC:", round(aic_robust, 2), ")") cat("Stepwise Model
- MAE:", round(mae_stepwise, 2), "| RMSE:", round(rmse_stepwise, 2), "| AIC:", round(aic_stepwise, 2), ")")
cat("BoxCox Model - MAE:", round(mae_boxcox, 2), "| RMSE:", round(rmse_boxcox, 2), "| AIC:",
round(aic_boxcox, 2), ")")
```

Preparing new observations

```
new_data <- data.frame( Age = log1p(c(45, 60)),
BMI = sqrt(c(28, 35)),
Insulin = log1p(c(90, 150)),
BloodPressure = sqrt(c(80, 90))
)
```

Prediction using stepwise model

```
predicted_transformed <- predict(model_stepwise, newdata = new_data)
```

Applying inverse Box-Cox transformation

Box-Cox lambda value is

```
lambda <- 0.06060606
inverse_boxcox <- function(y, lambda) { if (abs(lambda) < 1e-4) { return(exp(y))
} else { return((y * lambda + 1)^(1 / lambda))
}}
predicted_glucose <- inverse_boxcox(predicted_transformed, lambda)
```

Displaying results

```
result_df <- cbind( New_Patient = c("Patient 1", "Patient 2"), Predicted_Glucose_Transformed =
round(predicted_transformed, 4), Predicted_Glucose = round(predicted_glucose, 2) )
print(result_df)
```